# CPEN 502 Assignment-b: Reinforcement Learning (Look Up Table)

Ali Asgari Khoshouyeh (Student #24868739)

24. November 2021

## 1 Q Learning Robot

(2) Once you have your robot working, measure its learning performance as follows:

a) Draw a graph of a parameter that reflects a measure of progress of learning and comment on the convergence of learning of your robot.

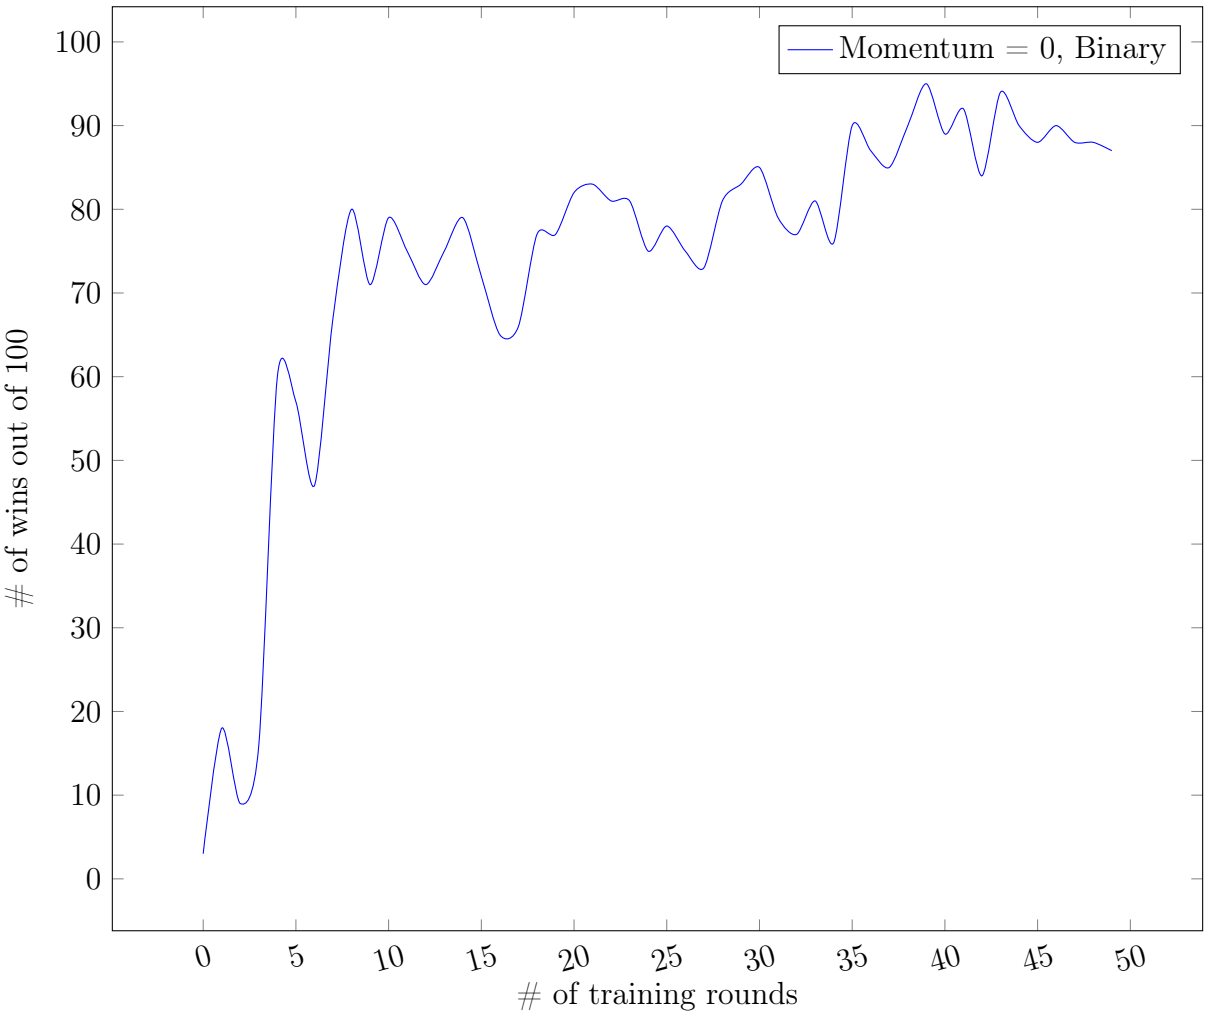


b) Using your robot, show a graph comparing the performance of your robot using on-policy learning vs off-policy learning.

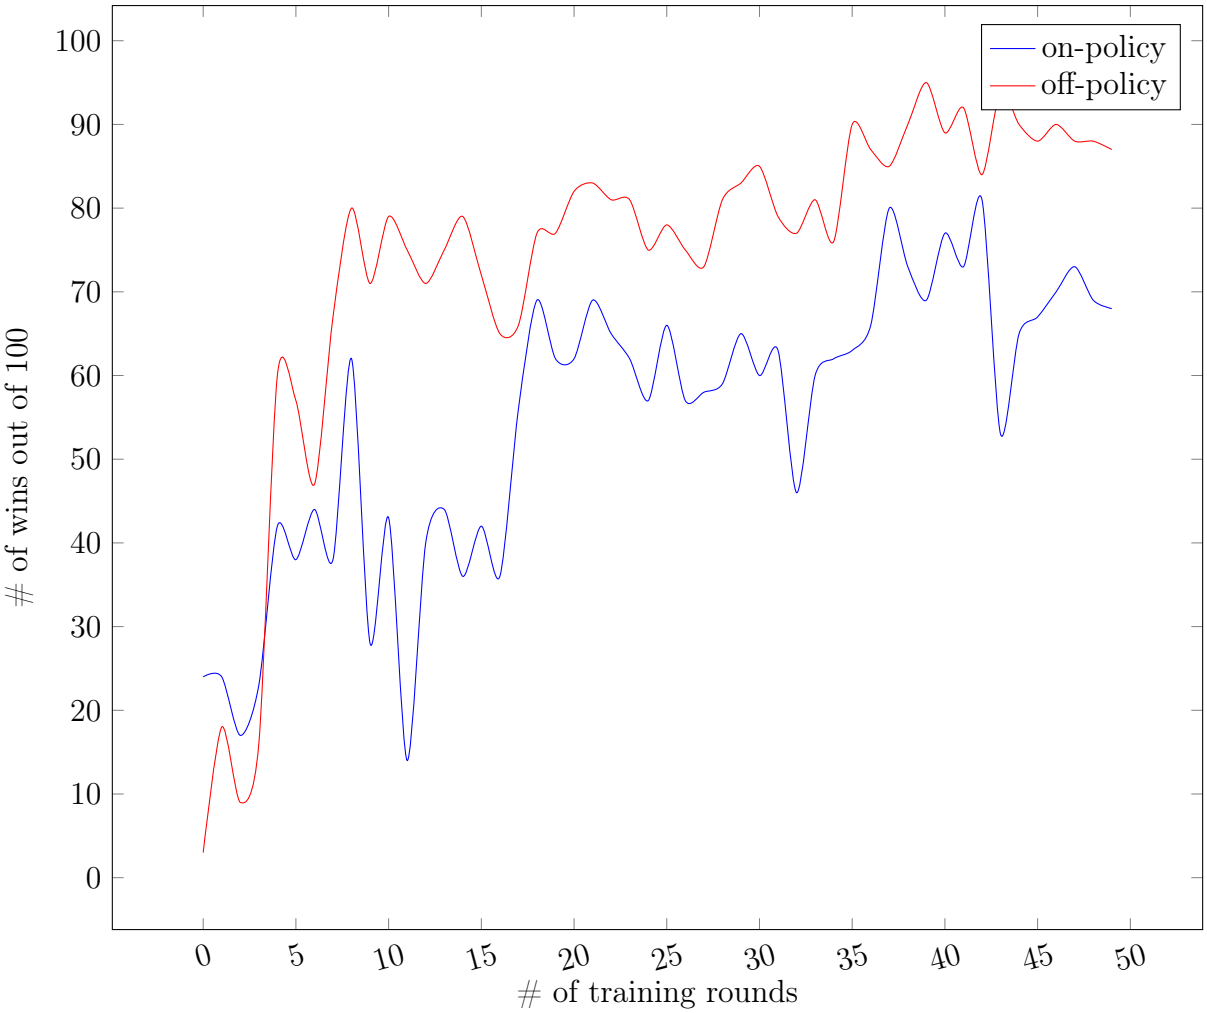


c) Implement a version of your robot that assumes only terminal rewards and show & compare its behaviour with one having intermediate rewards.

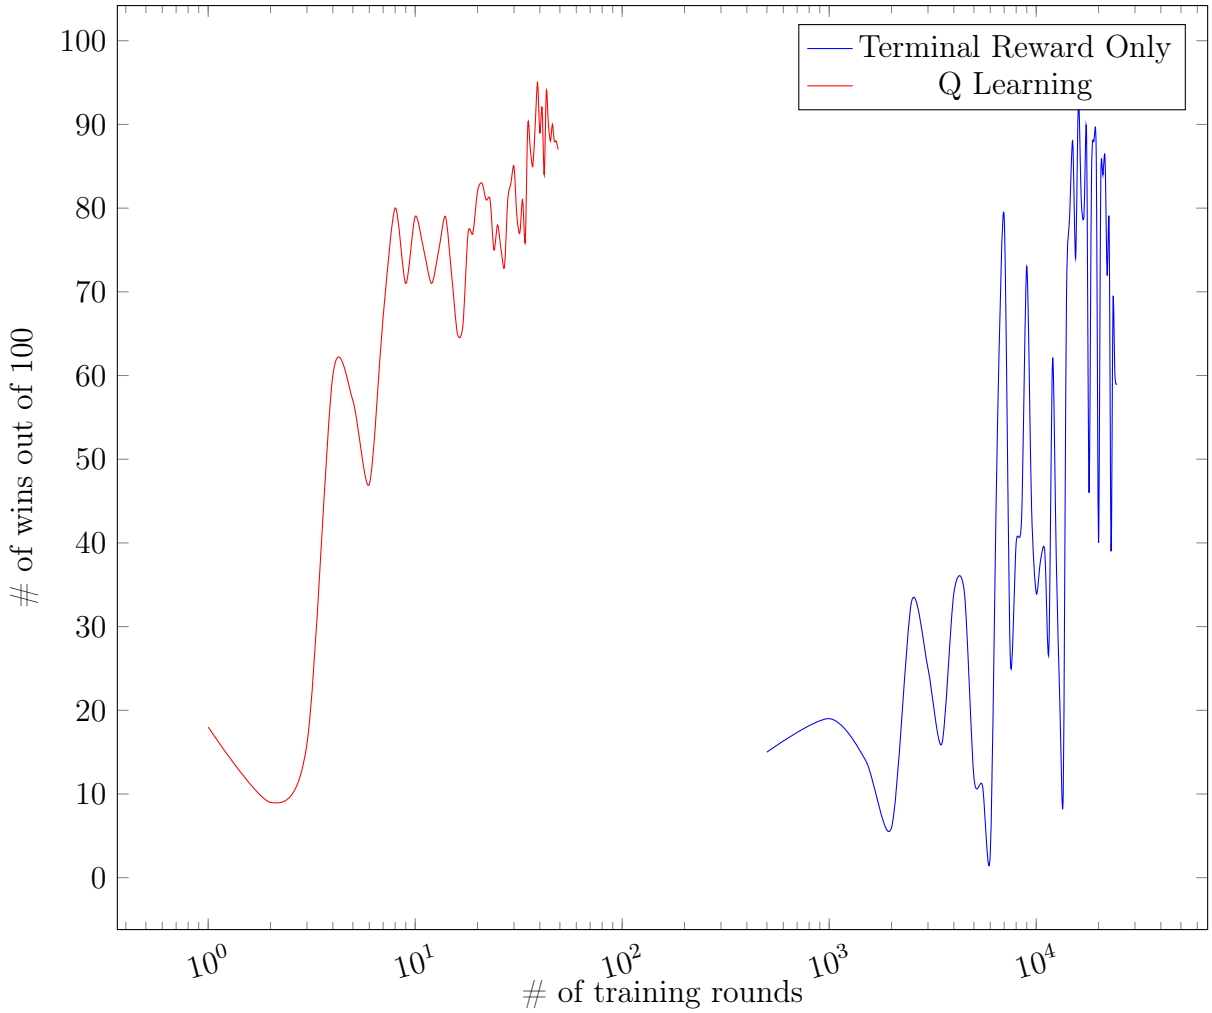


## 2 Role of $\epsilon$

(3) This part is about exploration. While training via RL, the next move is selected randomly with probability $\epsilon$ and greedily with probability $1 - \epsilon$

a) Compare training performance using different values of e including no exploration at all. Provide graphs of the measured performance of your tank vs $\epsilon$

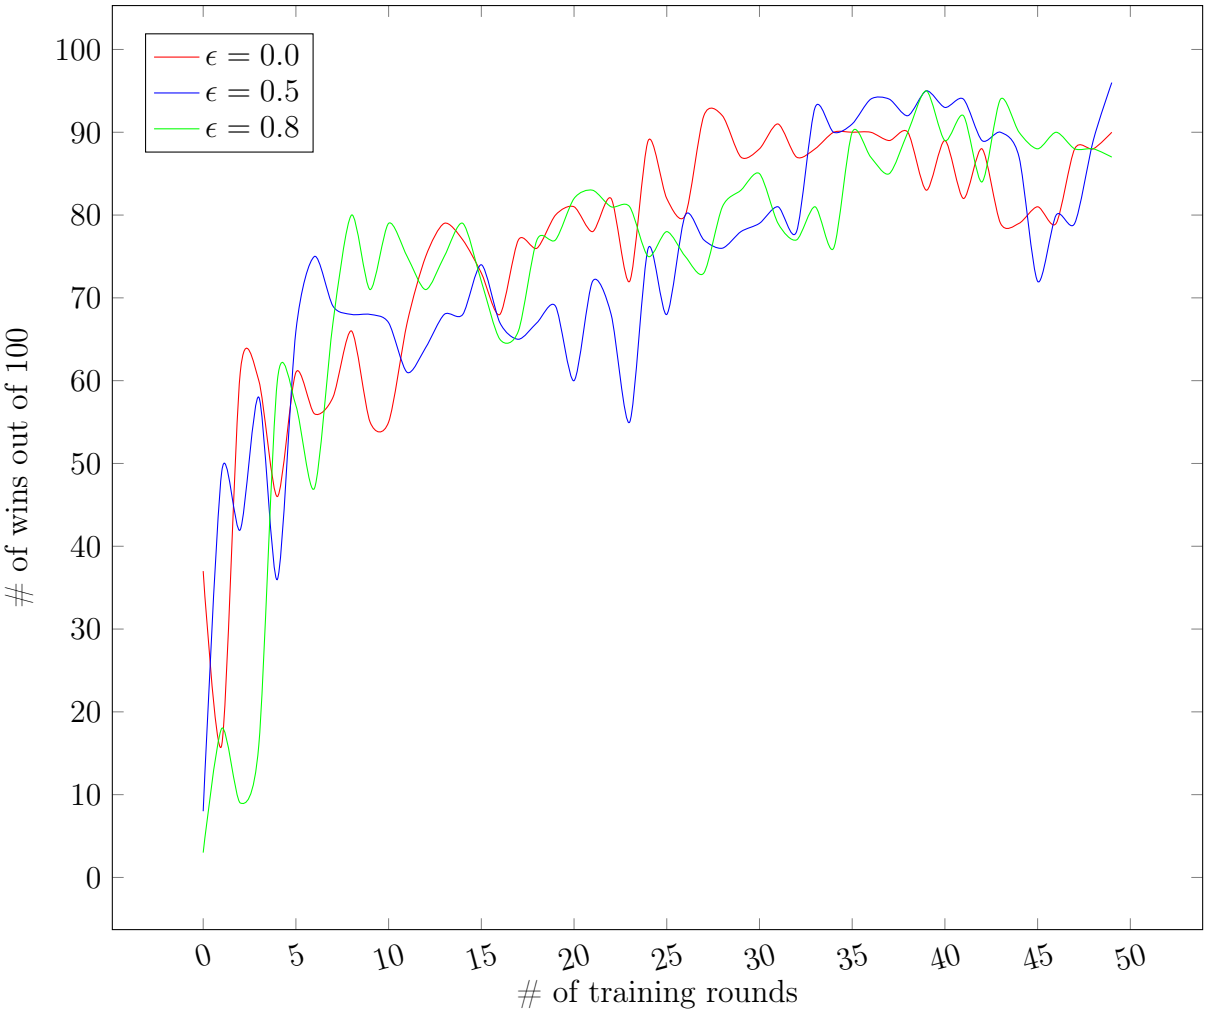
$\epsilon = 0.0$
$\epsilon = 0.5$
$\epsilon = 0.8$
# of wins out of 100
# of training rounds
0
10
20
30
40
50
60
70
80
90
100
0
5
10
15
20
25
30
35
40
45
50

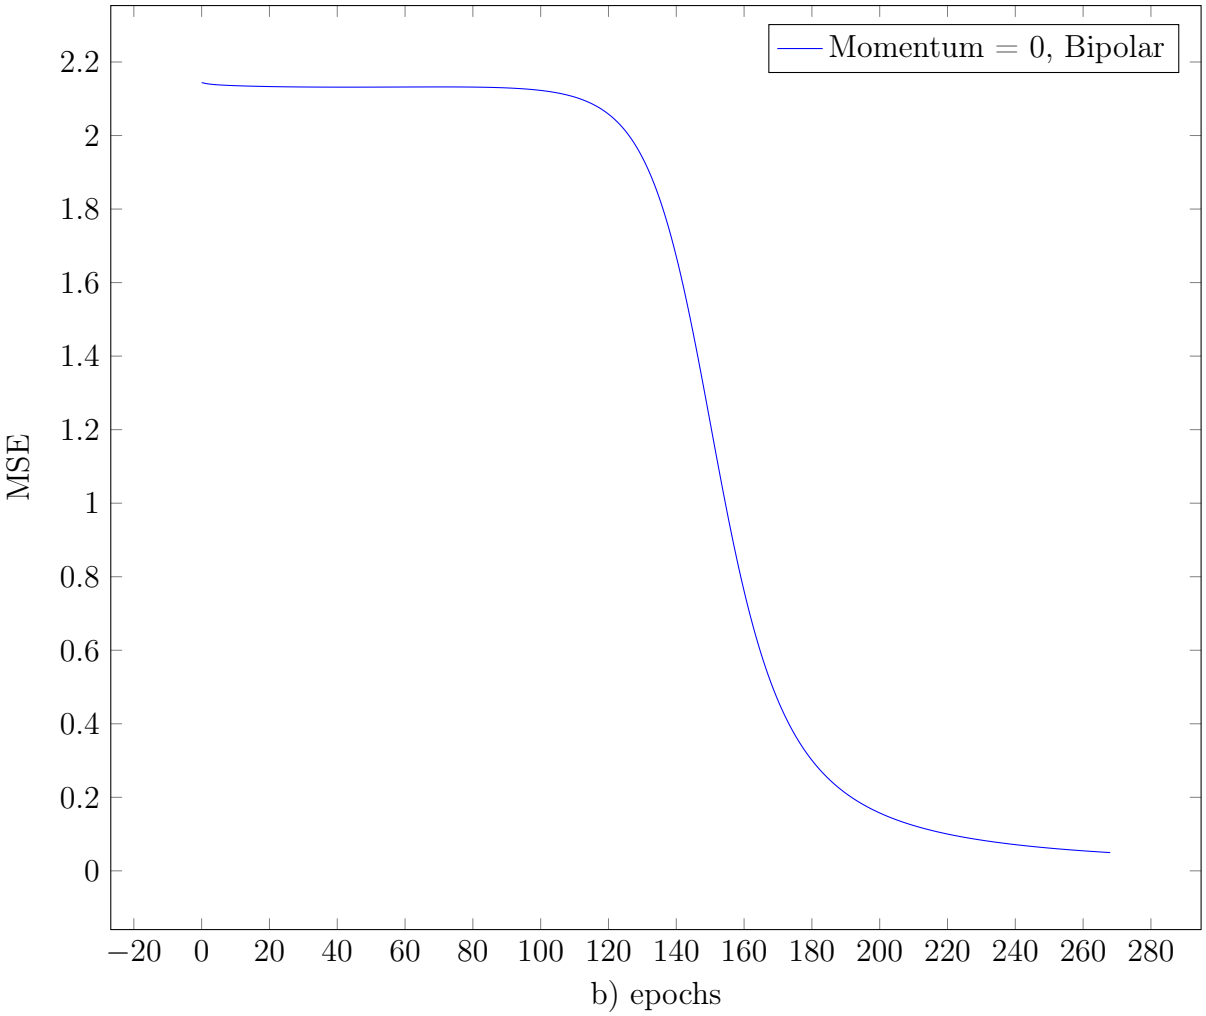


# Appendices

## A Source Codes

```
package autograd;

import jdk.jshell.spi.ExecutionControl;

public class Addition extends Operator {
    @Override
    public double evaluate(IVariable[] operands) {
        double result = 0.;
        for (IVariable operand :
                operands) {
            result += operand.evaluate();
        }
        return result;
    }

    @Override
```

```
17      public void backwards(IVariable[] operands, IVariable[] sources, double
            gradient) throws ExecutionControl.NotImplementedException {
18          for (IVariable o :
19                  operands) {
20              o.backward(sources, gradient);
21          }
22      }
23  }
```

Listing 1: autograd/Addition.java

```
1   package autograd;
2
3   import jdk.jshell.spi.ExecutionControl;
4
5   public class Exponentiation extends Operator {
6       @Override
7       public double evaluate(IVariable[] operands) {
8           if (operands.length != 2) {
9               throw new IllegalArgumentException("Exponentiation accepts 2
                    arguments.");
10          }
11          return Math.pow(operands[0].evaluate(), operands[1].evaluate());
12      }
13
14      @Override
15      public void backwards(IVariable[] operands, IVariable[] sources, double
            gradient) throws ExecutionControl.NotImplementedException {
16          IVariable baseVariable = operands[0];
17          var baseValue = baseVariable.evaluate();
18          IVariable exponentVariable = operands[1];
19          var exponentValue = exponentVariable.evaluate();
20          if (exponentVariable.getParameters().length > 1) {
21              throw new ExecutionControl.NotImplementedException("Back
                    propagation to the exponent is not implemented.");
22          }
23          var gradientToPropagate = Math.pow(gradient * baseValue *
                exponentValue, exponentValue - 1);
24          baseVariable.backward(sources, gradientToPropagate);
25      }
26  }
```

Listing 2: autograd/Exponentiation.java

```
1   package autograd;
2
3   public interface IInitializer {
4       double next();
5   }
```

Listing 3: autograd/IInitializer.java

```
1   package autograd;
2
3
4   import jdk.jshell.spi.ExecutionControl;
5
6   public interface IOperator {
```

```
    IVariable apply(IVariable... operands);

    double evaluate(IVariable[] operands);

    void backwards(IVariable[] operands, IVariable[] sources, double
        gradient) throws ExecutionControl.NotImplementedException;
}
```

Listing 4: autograd/IOperator.java

```
package autograd;

import jdk.jshell.spi.ExecutionControl;

public interface IVariable {
    double evaluate();

    void backward(IVariable[] sources, double gradient) throws
        ExecutionControl.NotImplementedException;

    Parameter[] getParameters();
}
```

Listing 5: autograd/IVariable.java

```
package autograd;

import jdk.jshell.spi.ExecutionControl;

public class Multiplication extends Operator {

    @Override
    public double evaluate(IVariable[] operands) {
        double result = 1.;
        for (IVariable operand :
                operands) {
            result *= operand.evaluate();
        }
        return result;
    }

    @Override
    public void backwards(IVariable[] operands, IVariable[] sources, double
         gradient) throws ExecutionControl.NotImplementedException {
        validateOperands(operands);
        var multiplier = operands[0];
        var multiplicand = operands[1];
        var multiplierValue = multiplier.evaluate();
        var multiplicandValue = multiplicand.evaluate();
        multiplier.backward(sources, gradient * multiplicandValue);
        multiplicand.backward(sources, gradient * multiplierValue);
    }
}
```

Listing 6: autograd/Multiplication.java

```
package autograd;

```

```java
import jdk.jshell.spi.ExecutionControl;

public class Negation extends Operator {

    public Negation() {
        this.numberOfOperands = 1;
    }

    @Override
    public double evaluate(IVariable[] operands) {
        validateOperands(operands);
        return -operands[0].evaluate();
    }

    @Override
    public void backwards(IVariable[] operands, IVariable[] sources, double
        gradient) throws ExecutionControl.NotImplementedException {
        operands[0].backward(sources, -gradient);
    }
}
```

Listing 7: autograd/Negation.java

```java
package autograd;

import jdk.jshell.spi.ExecutionControl;

import java.util.Arrays;
import java.util.HashSet;

public class Operation implements IVariable {
    private final IOperator operator;
    private final IVariable[] operands;

    public Operation(IOperator operator, IVariable... operands) {
        this.operator = operator;
        this.operands = operands;
    }

    @Override
    public double evaluate() {
        return operator.evaluate(operands);
    }

    @Override
    public void backward(IVariable[] sources, double gradient) throws
       ExecutionControl.NotImplementedException {
        operator.backwards(operands, sources, gradient);
    }


    @Override
    public Parameter[] getParameters() {
        HashSet<Parameter> result = new HashSet<>();
        for (IVariable o :
                this.operands) {
            result.addAll(Arrays.asList(o.getParameters()));
        }
```

```
        return result.toArray(new Parameter[0]);
    }

    public IVariable[] getOperands() {
        return operands;
    }
}
```

Listing 8: autograd/Operation.java

```
package autograd;

public abstract class Operator implements IOperator {
    protected Integer numberOfOperands;

    public Operator() {
        this.numberOfOperands = null;
    }

    @Override
    public IVariable apply(IVariable... operands) {
        return new Operation(this, operands);
    }

    protected void validateOperands(IVariable[] operands) {
        if (this.numberOfOperands == null) {
            return;
        }
        if (operands.length != this.numberOfOperands) {
            throw new IllegalArgumentException(String.format("%s accepts
                only one operand.", this.getClass().getName()));
        }
    }
}
```

Listing 9: autograd/Operator.java

```
package autograd;

import java.util.Arrays;

public class Parameter implements IVariable {
    private double value;
    private double gradient;
    private boolean trainable;
    private int layer;

    public Parameter() {

    }

    public Parameter(double value) {
        this.value = value;
        trainable = true;
    }

    public Parameter(double value, boolean trainable) {
        this.value = value;
```

```
22          this.trainable = trainable;
23      }
24
25      public static IVariable[] createTensor(double[] desired) {
26          var result = new Parameter[desired.length];
27          for (int i = 0; i < result.length; i++) {
28              result[i] = new Parameter(desired[i]);
29          }
30          return result;
31      }
32
33      @Override
34      public double evaluate() {
35          return value;
36      }
37
38      @Override
39      public void backward(IVariable[] sources, double gradient) {
40          if (Arrays.stream(sources).anyMatch(x -> x == this)) {
41              setGradient(gradient + getGradient());
42          }
43      }
44
45      @Override
46      public Parameter[] getParameters() {
47          return new Parameter[]{this};
48      }
49
50      public double getValue() {
51          return this.value;
52      }
53
54      public void setValue(double value) {
55          this.value = value;
56      }
57
58      public double getGradient() {
59          return gradient;
60      }
61
62      private void setGradient(double gradient) {
63          this.gradient = gradient;
64      }
65
66      public boolean isTrainable() {
67          return this.trainable;
68      }
69
70      public void zeroGradient() {
71          this.setGradient(0);
72      }
73
74      public int getLayer() {
75          return layer;
76      }
77
78      public void setLayer(int layer) {
79          this.layer = layer;
```

```
    }
}
```

Listing 10: autograd/Parameter.java

```
package autograd;

import jdk.jshell.spi.ExecutionControl;

public class Sigmoid extends Operator {

    public Sigmoid() {
        this.numberOfOperands = 1;
    }

    @Override
    public double evaluate(IVariable[] operands) {
        if (operands.length != 1) {
            throw new IllegalArgumentException("Sigmoid operator only
                accepts one operand");
        }
        return 1. / (1 + Math.exp(-operands[0].evaluate()));
    }

    @Override
    public void backwards(IVariable[] operands, IVariable[] sources, double
        gradient) throws ExecutionControl.NotImplementedException {
        validateOperands(operands);
        var x = operands[0];
        var y = evaluate(operands);
        x.backward(sources, gradient * y * (1 - y));
    }
}
```

Listing 11: autograd/Sigmoid.java

```
package autograd;

import java.util.Random;

public class UniformInitializer implements IInitializer {

    double a;
    double b;
    Random random;

    public UniformInitializer(double a, double b) {
        this.a = a;
        this.b = b;
        this.random = new Random();
    }

    @Override
    public double next() {
        return random.nextDouble() * (b - a) + a;
    }
}
```

Listing 12: autograd/UniformInitializer.java

```java
package dataset;

public class BinaryToBipolarWrapper implements IDataSet {

    IDataSet binaryDataSet;

    public BinaryToBipolarWrapper(IDataSet binaryDataSet) {
        this.binaryDataSet = binaryDataSet;
    }

    @Override
    public DataPoint next() {
        DataPoint result = binaryDataSet.next();
        if (result == null) return null;
        double[] x = result.getX().clone();
        double[] y = result.getY().clone();
        for (int i = 0; i < x.length; i++) {
            x[i] = 2 * x[i] - 1;
        }
        for (int i = 0; i < y.length; i++) {
            y[i] = 2 * y[i] - 1;
        }
        return new DataPoint(x, y);
    }

    @Override
    public void reset() {
        binaryDataSet.reset();
    }
}
```

Listing 13: dataset/BinaryToBipolarWrapper.java

```java
package dataset;

public class DataPoint {
    private final double[] x;
    private final double[] y;

    public DataPoint(double[] x, double[] y) {
        this.x = x;
        this.y = y;
    }

    public double[] getY() {
        return y;
    }

    public double[] getX() {
        return x;
    }
}
```

Listing 14: dataset/DataPoint.java

```java
package dataset;

public interface IDataSet {
```

```
    DataPoint next();

    void reset();
}
```

Listing 15: dataset/IDataSet.java

```
package dataset;

public class XORBinaryDataSet implements IDataSet {

    protected double[][] x;
    protected double[] y;
    private int index;

    public XORBinaryDataSet() {
        index = 0;
        x = new double[][]{
                {0., 0.},
                {0., 1.},
                {1., 0.},
                {1., 1.},
        };
        y = new double[]{
                0.,
                1.,
                1.,
                0.,
        };
    }

    @Override
    public DataPoint next() {
        if (index < x.length) {
            var result = new DataPoint(x[index], new double[]{y[index]});
            index++;
            return result;
        }
        return null;
    }

    @Override
    public void reset() {
        index = 0;
    }
}
```

Listing 16: dataset/XORBinaryDataSet.java

```
package nn;

import autograd.IVariable;
import autograd.Parameter;

public class BipolarSigmoid implements ILayer {

    @Override
    public IVariable[] apply(IVariable[] input) {
```

```
        var sigmoid = new autograd.Sigmoid();
        var scalar = new Parameter(2, false);
        var constant = new Parameter(-1, false);
        var addition = new autograd.Addition();
        var multiplication = new autograd.Multiplication();
        var result = new IVariable[input.length];
        for (int i = 0; i < input.length; i++) {
            result[i] = addition.apply(
                    multiplication.apply(
                            scalar,
                            sigmoid.apply(input[i])),
                    constant
            );
        }
        return result;
    }
}
```

Listing 17: nn/BipolarSigmoid.java

```
package nn;

import java.util.ArrayList;

public class ConvergenceCollector implements IFitCallback {
    ArrayList<Double> loss;

    public ConvergenceCollector() {
        this.loss = new ArrayList<>();
    }

    @Override
    public void collect(int epoch, double loss) {
        this.loss.add(loss);
    }

    public int getEpochs() {
        return loss.size();
    }

    @Override
    public String toString() {
        StringBuilder sb = new StringBuilder();
        for (int i = 0; i < loss.size(); i++) {
            sb.append(+i + " " + loss.get(i) + "\n");
        }
        return sb.toString();
    }
}
```

Listing 18: nn/ConvergenceCollector.java

```
package nn;

import autograd.IInitializer;
import autograd.IVariable;
import autograd.Parameter;
import autograd.UniformInitializer;
```

```
7
8  public class Factory {
9      public static Model createNeuralNetwork(int[] sizes, ILayer activation,
           IInitializer initializer) {
10         if (sizes.length < 2) {
11             throw new IllegalArgumentException("Sizes must at least contain
                   2 integers for the first and the second layer.");
12         }
13         var inputs = new Parameter[sizes[0]];
14         for (int i = 0; i < inputs.length; i++) {
15             inputs[i] = new Parameter(initializer.next());
16         }
17         IVariable[] lastLayerOutput = inputs;
18         for (int i = 1; i < sizes.length; i++) {
19             lastLayerOutput = new Linear(sizes[i - 1], sizes[i],
                   initializer).apply(lastLayerOutput);
20             lastLayerOutput = activation.apply(lastLayerOutput);
21         }
22         return new Model(inputs, lastLayerOutput);
23     }
24
25     public static Model createNeuralNetwork(int[] sizes, ILayer activation)
           {
26         return createNeuralNetwork(sizes, activation, new
               UniformInitializer(-0.5, 0.5));
27     }
28 }
```

Listing 19: nn/Factory.java

```
1  package nn;
2
3  public interface IFitCallback {
4      void collect(int epoch, double loss);
5  }
```

Listing 20: nn/IFitCallback.java

```
1  package nn;
2
3  import autograd.IVariable;
4
5  public interface ILayer {
6      IVariable[] apply(IVariable[] input);
7  }
```

Listing 21: nn/ILayer.java

```
1  package nn;
2
3  import autograd.*;
4
5  public class Linear implements ILayer {
6      private final IVariable[][] weight;
7      private final IVariable[] bias;
8
9      public Linear(int inFeatures, int outFeatures, IInitializer initializer
           ) {
```

```
        this.weight = new Parameter[outFeatures][inFeatures];
        this.bias = new Parameter[outFeatures];
        for (int i = 0; i < outFeatures; i++) {
            for (int j = 0; j < inFeatures; j++) {
                this.weight[i][j] = new Parameter(initializer.next());
            }
            this.bias[i] = new Parameter(initializer.next());
        }
    }

    @Override
    public IVariable[] apply(IVariable[] input) {
        var result = new IVariable[this.weight.length];
        for (int i = 0; i < this.weight.length; i++) {
            int inputSize = this.weight[i].length;
            IVariable[] muls = new IVariable[inputSize + 1];
            for (int j = 0; j < inputSize; j++) {
                muls[j] = new Multiplication().apply(this.weight[i][j],
                    input[j]);
            }
            muls[inputSize] = this.bias[i];
            result[i] = new Addition().apply(muls);
        }
        return result;
    }


    private int getWidth() {
        return this.weight.length;
    }
}
```

Listing 22: nn/Linear.java

```
package nn;

import autograd.*;
import jdk.jshell.spi.ExecutionControl;
import optimization.ILoss;

public class MinimumSquaredError implements IVariable, ILoss {

    private final IVariable operation;
    private final Parameter[] desired;

    public MinimumSquaredError(IVariable[] output) {
        var negation = new Negation();
        var addition = new Addition();
        var multiplication = new Multiplication();
        var exponentiation = new Exponentiation();
        Parameter two = new Parameter(2, false);
        Parameter half = new Parameter(0.5, false);
        int length = output.length;
        desired = new Parameter[output.length];
        var summationTerms = new IVariable[length];
        for (int i = 0; i < length; i++) {
            desired[i] = new Parameter();
            summationTerms[i] = exponentiation.apply(
```

```
                    addition.apply(output[i], negation.apply(desired[i])),
                    two
            );
        }
        this.operation = multiplication.apply(addition.apply(summationTerms
            ), half);
    }

    @Override
    public double evaluate() {
        return operation.evaluate();
    }

    @Override
    public void backward(IVariable[] sources, double gradient) throws
        ExecutionControl.NotImplementedException {
        operation.backward(sources, gradient);
    }

    @Override
    public Parameter[] getParameters() {
        return this.operation.getParameters();
    }

    @Override
    public void setDesired(double[] desired) {
        for (int i = 0; i < this.desired.length; i++) {
            this.desired[i].setValue(desired[i]);
        }
    }
}
```

Listing 23: nn/MinimumSquaredError.java

```
package nn;

import autograd.IVariable;
import autograd.Operation;
import autograd.Parameter;
import dataset.DataPoint;
import dataset.IDataSet;
import jdk.jshell.spi.ExecutionControl;
import optimization.ILoss;
import optimization.IOptimizer;

import java.util.Arrays;
import java.util.HashSet;
import java.util.List;
import java.util.Map;
import java.util.stream.Collectors;

public class Model {
    private final Parameter[] input;
    private final IVariable[] output;

    public Model(Parameter[] input, IVariable[] output) {
        this.input = input;
        this.output = output;
```

```
25      }
26
27
28      public double[] evaluate(double[] input) {
29          var result = new double[output.length];
30          for (int i = 0; i < input.length; i++) {
31              this.input[i].setValue(input[i]);
32          }
33          for (int i = 0; i < output.length; i++) {
34              result[i] = output[i].evaluate();
35          }
36          return result;
37      }
38
39      public Parameter[] getParameters() {
40          HashSet<Parameter> result = new HashSet<>();
41          for (IVariable o :
42                  this.output) {
43              result.addAll(Arrays.asList(o.getParameters()));
44          }
45          return result.toArray(new Parameter[0]);
46      }
47
48      public Parameter[] getTrainableParameters() {
49          var results = new HashSet<Parameter>();
50          for (Parameter p :
51                  getParameters()) {
52              if (p.isTrainable()) {
53                  results.add(p);
54              }
55          }
56          for (Parameter p : input) {
57              results.remove(p);
58          }
59
60          return results.toArray(new Parameter[0]);
61      }
62
63      public IVariable[] getOutput() {
64          return output;
65      }
66
67      public double fit(IDataSet dataSet, IOptimizer optimizer, ILoss loss,
            int epochs, double lossLimit) throws ExecutionControl.
            NotImplementedException {
68          return fit(dataSet, optimizer, loss, epochs, lossLimit, (epoch, l)
                -> {
69          });
70      }
71
72      public double fit(IDataSet dataSet, IOptimizer optimizer, ILoss loss,
            int epochs, double lossLimit, IFitCallback callback) throws
            ExecutionControl.NotImplementedException {
73          var parameters = getTrainableParameters();
74          Map<Integer, List<Parameter>> layeredParameters = layerParameters(
                parameters);
75          if (epochs < 1) {
76              throw new IllegalArgumentException("At least one epochs
```

```
                required.");
77          }
78          double totalLoss = 0;
79          for (int i = 0; i < epochs; i++) {
80              totalLoss = 0;
81              dataSet.reset();
82              DataPoint dataPoint;
83              while ((dataPoint = dataSet.next()) != null) {
84                  setInput(dataPoint.getX());
85                  loss.setDesired(dataPoint.getY());
86                  totalLoss += loss.evaluate();
87                  for (Integer j : layeredParameters.keySet().stream().sorted
                        ().collect(Collectors.toList())) {
88                      Parameter[] layerParameters = layeredParameters.get(j).
                            toArray(new Parameter[0]);
89                      loss.backward(layerParameters, 1.);
90                      optimizer.update(layerParameters);
91                  }
92              }
93              callback.collect(i, totalLoss);
94              if (totalLoss < lossLimit) {
95                  break;
96              }
97          }
98          return totalLoss;
99      }
100
101     private Map<Integer, List<Parameter>> layerParameters(Parameter[]
            parameters) {
102         setLayers(getOutput(), 0);
103         return Arrays.stream(parameters).collect(Collectors.groupingBy(
                Parameter::getLayer));
104
105     }
106
107     private void setLayers(IVariable[] outputs, int layer) {
108         if (outputs.length == 0) return;
109         HashSet<IVariable> nextOutput = new HashSet<>();
110         for (IVariable i : outputs) {
111             if (i instanceof Parameter) {
112                 ((Parameter) i).setLayer(layer);
113             }
114             if (i instanceof Operation) {
115                 nextOutput.addAll(Arrays.asList(((Operation) i).getOperands
                        ()));
116             }
117         }
118         setLayers(nextOutput.toArray(new IVariable[0]), layer + 1);
119     }
120
121     private void setInput(double[] x) {
122         for (int i = 0; i < input.length; i++) {
123             input[i].setValue(x[i]);
124         }
125     }
126 }
```

Listing 24: nn/Model.java

```java
package nn;

import autograd.IVariable;

public class Sigmoid implements ILayer {

    @Override
    public IVariable[] apply(IVariable[] input) {
        var operator = new autograd.Sigmoid();
        var result = new IVariable[input.length];
        for (int i = 0; i < input.length; i++) {
            result[i] = operator.apply(input[i]);
        }
        return result;
    }
}
```

Listing 25: nn/Sigmoid.java

```java
package optimization;

import autograd.Parameter;

import java.util.HashMap;

public class GradientDescent implements IOptimizer {

    private final HashMap<Parameter, Double> lastDelta;
    private final double learningRate;
    private final double momentum;

    public GradientDescent(double learningRate, double momentum) {
        this.lastDelta = new HashMap<>();
        this.learningRate = learningRate;
        this.momentum = momentum;
    }

    @Override
    public void update(Parameter[] parameters) {
        for (Parameter p :
                parameters) {
            double delta = -p.getGradient() * learningRate + momentum *
                lastDelta.getOrDefault(p, 0.);
            p.setValue(p.getValue() + delta);
            p.zeroGradient();
            lastDelta.put(p, delta);
        }
    }
}
```

Listing 26: optimization/GradientDescent.java

```java
package optimization;

import autograd.IVariable;

public interface ILoss extends IVariable {
    void setDesired(double[] desired);
```

```
}
```

Listing 27: optimization/ILoss.java

```
package optimization;

import autograd.Parameter;

public interface IOptimizer {
    void update(Parameter[] parameters);
}
```

Listing 28: optimization/IOptimizer.java

```
package autograd;


import org.junit.Assert;
import org.junit.Test;

public class VariableTest {

    @Test
    public void testAddition() {
        Assert.assertEquals(new Addition().apply(new Parameter(12), new
            Parameter(2.)).evaluate(), 14., 0.);
    }

    @Test
    public void testVariableEvaluation() {
        Assert.assertEquals(new Parameter(250).evaluate(), 250., 0);
    }

}
```

Listing 29: autograd/VariableTest.java

```
package nn;

import autograd.Parameter;
import jdk.jshell.spi.ExecutionControl;
import org.junit.Assert;
import org.junit.Test;

import java.util.Arrays;

public class NeuralNetworkTest {

    @Test
    public void testNeuralNetworkFactory() {
        var model = Factory.createNeuralNetwork(new int[]{2, 4, 1}, new
            Sigmoid());
        var result = model.evaluate(new double[]{0, 0});
        Assert.assertEquals(result.length, 1);
    }

    @Test
    public void testNeuralNetworkGradient() throws ExecutionControl.
        NotImplementedException {
```

```
21          var model = Factory.createNeuralNetwork(new int[]{2, 4, 1}, new
                Sigmoid());
22          Parameter[] parameters = model.getTrainableParameters();
23          for (Parameter parameter :
24                  parameters) {
25              parameter.setValue(1);
26          }
27          var result = model.evaluate(new double[]{1, 0});
28          double[] desired = new double[]{1};
29          Assert.assertEquals(result.length, 1);
30          double delta = 1e-5;
31          double expected = 0.9892621636390686; // obtained by pytorch
32          Assert.assertEquals(result[0], expected, delta);
33          var loss = new MinimumSquaredError(model.getOutput());
34          loss.setDesired(desired);
35          loss.backward(parameters, 1);
36          var gradients = new double[parameters.length];
37          for (int i = 0; i < parameters.length; i++) {
38              gradients[i] = parameters[i].getGradient();
39          }
40
41          Assert.assertArrayEquals(Arrays.stream(new double[]{ // calculated
                by pytorch
42                  -0.000114063048386015, -0.00010046639363281429,
                        -0.00010046639363281429, -0.00010046639363281429,
                        -0.00010046639363281429, -1.197589335788507e-05,
                        -1.197589335788507e-05, -1.197589335788507e-05,
                        -1.197589335788507e-05, -1.197589335788507e-05,
                        -1.197589335788507e-05, -1.197589335788507e-05,
                        -1.197589335788507e-05, -0.0, -0.0, -0.0, -0.0
43          }).sorted().toArray(), Arrays.stream(gradients).sorted().toArray(),
                delta);
44      }
45
46      @Test
47      public void testNeuralNetworkGradientBipolar() throws ExecutionControl.
            NotImplementedException {
48          var model = Factory.createNeuralNetwork(new int[]{2, 4, 1}, new
                BipolarSigmoid());
49          Parameter[] parameters = model.getTrainableParameters();
50          for (Parameter parameter :
51                  parameters) {
52              parameter.setValue(1);
53          }
54          var result = model.evaluate(new double[]{1, -1});
55          double[] desired = new double[]{1};
56          Assert.assertEquals(result.length, 1);
57          double delta = 1e-5;
58          double expected = 0.8904789686203003; // obtained by pytorch
59          Assert.assertEquals(expected, result[0], delta);
60          var loss = new MinimumSquaredError(model.getOutput());
61          loss.setDesired(desired);
62          loss.backward(parameters, 1);
63          var gradients = new double[parameters.length];
64          for (int i = 0; i < parameters.length; i++) {
65              gradients[i] = parameters[i].getGradient();
66          }
67
```

```
68          Assert.assertArrayEquals(Arrays.stream(new double[]{ // calculated
                by pytorch
69                  -0.011338012292981148, -0.005239490419626236,
                        -0.005239490419626236, -0.005239490419626236,
                        -0.005239490419626236, -0.004458376672118902,
                        -0.004458376672118902, -0.004458376672118902,
                        -0.004458376672118902, -0.004458376672118902,
                        -0.004458376672118902, -0.004458376672118902,
                        -0.004458376672118902, 0.004458376672118902,
                        0.004458376672118902, 0.004458376672118902,
                        0.004458376672118902
70          }).sorted().toArray(), Arrays.stream(gradients).sorted().toArray(),
                delta);
71      }
72  }
```

Listing 30: nn/NeuralNetworkTest.java

```
1   package optimization;
2
3   import autograd.UniformInitializer;
4   import dataset.BinaryToBipolarWrapper;
5   import dataset.XORBinaryDataSet;
6   import jdk.jshell.spi.ExecutionControl;
7   import nn.*;
8   import org.junit.Assert;
9   import org.junit.Ignore;
10  import org.junit.Test;
11
12  import java.io.FileWriter;
13  import java.io.IOException;
14  import java.util.ArrayList;
15  import java.util.Comparator;
16  import java.util.Optional;
17
18  public class GradientDescentTest {
19
20      private final static int trials = 300;
21
22      @Ignore
23      @Test
24      public void TestSimpleGD() throws ExecutionControl.
            NotImplementedException {
25          var model = Factory.createNeuralNetwork(
26                  new int[]{2, 4, 1},
27                  new Sigmoid(),
28                  new UniformInitializer(-0.5, 0.5)
29          );
30          var dataSet = new XORBinaryDataSet();
31          var optimizer = new GradientDescent(0.2, 0.);
32          var loss = new MinimumSquaredError(model.getOutput());
33          double finalLoss = model.fit(dataSet, optimizer, loss, 40000, 0.05)
                ;
34          Assert.assertTrue("Big loss " + finalLoss, finalLoss < 0.05);
35      }
36
37      @Ignore("Skipping slow convergence tests.")
38      @Test
```

```
39      public void TestConvergence() throws ExecutionControl.
            NotImplementedException, IOException {
40          int diverged = 0;
41          ArrayList<ConvergenceCollector> stats = new ArrayList<>();
42          for (int i = 0; i < GradientDescentTest.trials; i++) {
43              var model = Factory.createNeuralNetwork(
44                      new int[]{2, 4, 1},
45                      new Sigmoid(),
46                      new UniformInitializer(-0.5, 0.5)
47              );
48              var dataSet = new XORBinaryDataSet();
49              var optimizer = new GradientDescent(0.2, 0.);
50              var loss = new MinimumSquaredError(model.getOutput());
51              var collector = new ConvergenceCollector();
52              double finalLoss = model.fit(dataSet, optimizer, loss, 40000,
                    0.05, collector);
53              stats.add(collector);
54              if (finalLoss > 0.05) {
55                  diverged += 1;
56              }
57          }
58          outputGraphData("a", stats);
59          Assert.assertTrue("Convergence with high probability busted!",
                diverged < 6);
60      }
61
62      private void outputGraphData(String assignmentPart, ArrayList<
            ConvergenceCollector> stats) throws IOException {
63          FileWriter of = new FileWriter("doc/" + assignmentPart + "_avg.tex"
                );
64          double average = stats.stream().mapToInt(ConvergenceCollector::
                getEpochs).average().getAsDouble();
65          of.write(String.valueOf(average));
66          of.close();
67
68          Optional<ConvergenceCollector> representative = stats.stream().min(
                Comparator.comparingDouble(c -> Math.abs(c.getEpochs() - average
                )));
69          of = new FileWriter("doc/" + assignmentPart + ".tex");
70          of.write(representative.get().toString());
71          of.close();
72      }
73
74      @Ignore("Skipping slow convergence tests.")
75      @Test
76      public void TestBipolarGD() throws ExecutionControl.
            NotImplementedException, IOException {
77          int diverged = 0;
78          int trials = GradientDescentTest.trials;
79          ArrayList<ConvergenceCollector> stats = new ArrayList<>();
80          for (int i = 0; i < trials; i++) {
81              var model = Factory.createNeuralNetwork(
82                      new int[]{2, 4, 1},
83                      new BipolarSigmoid(),
84                      new UniformInitializer(-0.5, 0.5)
85              );
86              var dataSet = new BinaryToBipolarWrapper(new XORBinaryDataSet()
                    );
```

```
87              var optimizer = new GradientDescent(0.2, 0.);
88              var loss = new MinimumSquaredError(model.getOutput());
89              var collector = new ConvergenceCollector();
90              double finalLoss = model.fit(dataSet, optimizer, loss, 3500,
                    0.05, collector);
91              if (finalLoss > 0.05) {
92                  diverged += 1;
93              }
94              stats.add(collector);
95          }
96          outputGraphData("b", stats);
97          Assert.assertTrue("Convergence with high probability busted! " +
                diverged + " failure out of " + trials, diverged < 6);
98      }
99
100     @Ignore
101     @Test
102     public void TestBipolarMomentumGD() throws ExecutionControl.
            NotImplementedException, IOException {
103
104         int diverged = 0;
105         int trials = GradientDescentTest.trials;
106         ArrayList<ConvergenceCollector> stats = new ArrayList<>();
107         for (int i = 0; i < trials; i++) {
108             var model = Factory.createNeuralNetwork(
109                     new int[]{2, 4, 1},
110                     new BipolarSigmoid(),
111                     new UniformInitializer(-0.5, 0.5)
112             );
113             var dataSet = new BinaryToBipolarWrapper(new XORBinaryDataSet()
                    );
114             var optimizer = new GradientDescent(0.2, 0.9);
115             var loss = new MinimumSquaredError(model.getOutput());
116             var collector = new ConvergenceCollector();
117             double finalLoss = model.fit(dataSet, optimizer, loss, 1000,
                    0.05, collector);
118             if (finalLoss > 0.05) {
119                 diverged += 1;
120             }
121             stats.add(collector);
122         }
123         outputGraphData("c", stats);
124         Assert.assertTrue("Convergence with high probability busted! " +
                diverged + " failure out of " + trials, diverged < 6);
125     }
126 }
```

Listing 31: optimization/GradientDescentTest.java